


Low-power Wireless Networking for the Internet of Things

Lab4: Data serialization for radio communication


Matteo Trobinger (matteo.trobinger@unitn.it)

Credits for some slides to:

Timofei Istomin, Pablo Corbalán, Enrico Soprana, Davide Vecchia



Struct type definition for the ping-pong message



```
typedef struct ping_pong_msg {  
    uint16_t sequence_number;  
    int16_t noise_floor;  
}  
ping_pong_msg_t;
```

```
typedef struct ping_pong_msg {  
    uint16_t sequence_number;  
    int16_t noise_floor;  
}  
attribute ((packed))  
ping_pong_msg_t;
```

What's that? A hack to avoid implementing a **serialization** framework

Serialization: the process of converting an object into a format that can be readily transported

Problems of C structures

There are three problems of C structures preventing their direct use in network messages:

1. **Byte order** — There are big-endian and little-endian platforms
2. **Unaligned access** — CPU can directly operate with a multi-byte number *only* if it is placed at the word boundary in memory
3. **Structure padding** — Compiler inserts holes between structure fields to optimize access. We don't want to send holes (and they are architecture-dependent)

Byte order

Least significant byte (LSB)

Let's take number **0x89ABCDEF** (hexadecimal).

In memory the number might be stored differently:

addr+0	EF
addr+1	CD
addr+2	AB
addr+3	89

LSB first

Little-endian
(e.g. x86)

MSB first

Big-endian
(some ARMs)

addr+0	89
addr+1	AB
addr+2	CD
addr+3	EF

In a network, computers **have to agree!**



- Numbers must be converted into **network byte order**, which is **big endian** [RFC 1700].

Otherwise, e.g., a **big-endian sensor** says the temperature is **100°C** and an **x86 server** thinks it's **1 677 721 600°C**.

Memory alignment

Most CPUs can work with multi-byte numbers **only** when they are placed at the **word boundary**

- 32-bit CPUs have words of 4 bytes



Unaligned access

```
uint16_t val;  
memcpy(&val, packetbuf_dataptr(), 2);  
printf("%d", val);
```



OK!

```
uint16_t* val = (uint16_t*)packetbuf_dataptr();  
printf("%d", *val);
```



Error!

- When you **receive anything, copy it first to a variable, then use!**
- To **send a value, copy it to the payload from a variable!**

Structure padding

Because of the memory alignment requirement, the compiler sometimes has to insert holes in structures

```
struct {  
    uint8_t  a;  
    uint32_t b;  
    uint16_t c;  
}
```

Size 7?

32-bit CPU →

```
struct {  
    uint8_t  a; // 3-byte hole  
    uint32_t b;  
    uint16_t c; // 2-byte hole  
}
```

Size 12

Packed structs

There's a way to tell the compiler to **disable padding** for a structure

```
struct {  
    uint8_t a;  
    // 3-byte hole  
    uint32_t b;  
    uint16_t c;  
    // 2-byte hole  
}  
__attribute__((packed))
```



Is it enough? It might be, if you care more about *simplicity* rather than *efficiency* and *portability*.

Why packed structs are not that good

- **They do not solve the endianness problem**
 - It's the programmer's job to take care of it
- **The code that the compiler generates to work with packed structs is less efficient**
 - Because the struct members are not properly aligned

Packed structs are inefficient

```
struct foo {int a;} x;
```

```
x.a=0x1234;
```

Compiled to

```
mov #0x1234, @x
```

```
struct bar{int a;}
```

```
__attribute__((packed)) y;
```

```
y.a=0x1234;
```

Compiled to

```
mov #y, r15  
mov.b #0x34, @r15  
mov.b #0x12, 1(r15)
```

Serialization (marshaling)


The correct way: for every *structure* to be sent over the network, define 2 functions

1. To *serialize* the structure to a byte buffer
2. To *deserialize* a byte buffer into the structure

Example: assume you need to send/receive a message containing the following 3 fields:


- **a:** 1 byte, **b:** 2 bytes, **c:** 4 bytes

First thing to do: define (*document*) the message format, e.g.:



Specify also the network byte order! We'll use **big-endian** here (and our platform might be little-endian)

Serialization



Values to TX:
a = 0x01
b = 0x2345
c = 0x6789ABCD


```
//pl is a pointer to the  
//payload portion of the  
//packetbuf (uint8_t*)
```

```
pl[0] = a;  
  
pl[1] = b >> 8;  
pl[2] = b;  
  
pl[3] = c >> 24;  
pl[4] = c >> 16;  
pl[5] = c >> 8;  
pl[6] = c;
```


pl[.] is a 8-bit
variable!

```
//pl is a pointer to the  
//payload portion of the  
//packetbuf (uint8_t*)
```


```
pl[0] = 0x01;  
  
pl[1] = 0x0023;  
pl[2] = 0x2345;  
  
pl[3] = 0x00000067;  
pl[4] = 0x00006789;  
pl[5] = 0x006789AB;  
pl[6] = 0x6789ABCD;
```



Deserialization



```
//pl is a pointer to the  
//received payload  
//(uint8_t*)  
  
a = pl[0];  
  
b = (uint16_t) pl[1] << 8;  
b |= (uint16_t) pl[2];  
  
c = (uint32_t) pl[3] << 24;  
c |= (uint32_t) pl[4] << 16;  
c |= (uint32_t) pl[5] << 8;  
c |= (uint32_t) pl[6];
```



```
//pl is a pointer to the  
//received payload  
//(uint8_t*)  
  
a = 01;  
  
b = 0x2300;  
b |= 0x0045; } After bitwise OR  
b = 0x2345  
  
c = 0x67000000;  
c |= 0x00890000; } After all  
c |= 0x0000AB00; } bitwise OR  
c = 0x6789ABCD  
c |= 0x000000CD;
```

Serialization (marshaling)

Instead of hard-coding indexes, it is convenient to use a **progressing counter** (`idx`)

Serialize

```
uint16_t idx = 0;  
  
pl[idx++] = a;  
  
pl[idx++] = b >> 8;  
pl[idx++] = b;  
  
pl[idx++] = c >> 24;  
pl[idx++] = c >> 16;  
pl[idx++] = c >> 8;  
pl[idx++] = c;
```

Deserialize

```
uint16_t idx = 0;  
  
a = pl[idx++];  
  
b = (uint16_t)pl[idx++] << 8;  
b |= (uint16_t)pl[idx++];  
  
c = (uint32_t)pl[idx++] << 24;  
c |= (uint32_t)pl[idx++] << 16;  
c |= (uint32_t)pl[idx++] << 8;  
c |= (uint32_t)pl[idx++];
```

Alert: Request access to CLOVES!

How to get access:

Fill in the form on Moodle with your unitn.it account to obtain your testbed credentials.

Deadline:

Registration **must be completed by October 26th at 23:59**



Code Templates

Download and unzip the provided code

- \$ unzip Lab4-exercise.zip

Go to the code directory

- \$ cd Lab4-exercise/chain-template

To compile:

- \$ make or make TARGET=SKY

When everything seems to work, test it in Cooja:

- \$ cooja chain.csc &

Discover who's around, forward messages to neighbors!

Neighbor discovery:

- Periodically send beacons (*in broadcast*) to announce the node's presence
- Upon receiving a beacon, store the sender address in a neighbor table


Message forwarding:

- On button press, send (*in unicast*) a packet to a random neighbor (the packet format is specified in the next slide!)
- When a unicast packet arrives, increment the hop count and forward the packet to a random neighbor

The template implements everything **except for message (de)serialization**

Finished? Try to improve the current protocol!
If you want, discuss your ideas/solution with me 😊

Unicast message format



Network byte order: big-endian