


# Low-power Wireless Networking for the Internet of Things

**Lab (6 and) 7:**  
**Data collection with many-to-one routing**

Matteo Trobinger ([matteo.trobinger@unitn.it](mailto:matteo.trobinger@unitn.it))

Credits for some slides to:

Pablo Corbalán, Timofei Istomin




# Labs 6-7 Goal

Implement a many-to-one routing-based data collection protocol capable of collecting sensor readings from any node of a multi-hop wireless network towards a single intended destination, called the sink.


## **Step 1 — Lab 6:**

Acquire topology information: Building a routing tree  
→ Routing metric: Hop count



## **Step 2 — Lab 7:**

Leverage such routes to reliably forward data packets across the network, from sources to the sink node



# Lab 6 Recap:

## How to build a routing tree for an unknown topology


# Lab 6 – Building the three (routing)

**S0:** Initially—when the application calls `my_collect_open()`—we have a *disconnected* network.

We need to:

1. Initialize our connection object
2. Open a connection with the underlying RIME primitives (broadcast and unicast)
3. Start the procedure to construct our routing tree **for the first time**

→ We need to build a routing tree with **the sink** (node 1) as the **root** (data collection point).




# Lab 6 – Building the tree (routing)

**S1:** The sink sends a beacon message in broadcast with **seqn = 0** and metric **h = 0**.

**S2:** If RSSI of the beacon > threshold, nodes 2 and 3 compare **h, seqn** of the received message against their current metrics and join the network.

- Select node 1 (the sink) as their **parent**.
- Set their own metrics to **h=1, seqn=0**.
- Prepare themselves to send, after a **small random delay**, a new beacon message in broadcast with **h=1, seqn=0**.




# Lab 6 – Building the tree (routing)

**S3:** Node 3 sends a broadcast message with **seqn = 0** and metric **h = 1**.

Node 4 (i) checks RSSI > threshold, (ii) joins the network with metrics **h = 2**, **seqn = 0**, (iii) selects node 3 as the parent, and (iv) prepares itself to TX after a **small, random delay** a beacon message in broadcast.

Nodes 1, 2 **do not** update their routing information as their metrics **do not improve** by selecting node 3 as parent.




# Lab 6 – Building the tree (routing)

**S4:** Node 2 sends a broadcast message with **seqn = 0** and metric **h = 1**.


Node 5 (i) checks RSSI > threshold, (ii) joins the network with metrics **h = 2, seqn = 0**, (iii) selects node 2 as the parent, and (iv) prepares itself to TX after a **small, random delay** a beacon message in broadcast.

Nodes 1, 3, 4 **do not** update their routing information as their metrics **do not improve** by selecting node 2 as parent.




# Lab 6 – Building the tree (routing)

**S5:** Node 4 sends a broadcast message with **seqn = 0** and metric **h = 2**.



**S6:** Node 5 sends a broadcast message with **seqn = 0** and metric **h = 2**.



In both cases **no routing changes** occur in the network!

# Lab 6 – Periodic tree updates

**Afterwards:** Refresh routes *periodically* by repeating the **whole process** from scratch.


**Why?** To cope with potential network changes that could *drastically modify* our previously learned topology.

**How?**

**Sink:** *Periodically increases the beacon sequence number* and starts a *new beacon flood* to (re)build the routing topology

**Nodes:** Upon receiving a **new** beacon message (i.e., seqn > highest seqn ever seen)

- Check RSSI > threshold,
- Accept the new metric (hop count, parent, ...) without checking it against the old one, and
- After a small random delay TX the updated beacon message in broadcast




# Lab 7

## Let's start forwarding!

Hey App!

# Implementing forwarding

When the routing is done,  
implement the forwarding mechanism!



You just need to complete two functions in `my_collect.c`:

1. `my_collect_send()` — called periodically *on non-sink nodes* by the application (`app.c`) to start sending data towards the sink
  - *Prepare* the data packet
  - Send the packet to the source's parent in *unicast*
2. `uc_recv()` — called by the underlying unicast layer to inform the `my_collect` layer when a data packet arrives
  - Forwarder (non-sink node): relay the received packet *to the node's parent* (again in unicast) to ensure forwarding progress
  - Sink: inform the application that a new data packet has been received (i.e., call `recv_cb`, the `recv` callback of `my_collect`)



# Forwarding at the data source – TODO 5

The application calls `my_collect_send()` to send data to the sink. This function should:

- Return **-1** (failure) if the node **is not yet connected** to the network
- Otherwise,
  1. *Prepare and attach a protocol header to the data packet;*
  2. Send the data packet to the source's parent in unicast
  3. Return the status returned by the unicast send function




## Our protocol header:

```
struct collect_header {  
    linkaddr_t source;  
    uint8_t hops;  
} __attribute__((packed));
```

Originator (source) node address: set at the source node, **never modified** along the path

Hop count, set to 0 at the source, **incremented every time** the packet is **received**



# Adding a header (at the source) – TODO 5


When the application sends a packet, it (i) **fills the packet buffer** with its message **starting from** `packetbuf_dataptr()`, and (ii) calls `my_collect_send()`

Packet buffer



The collect layer needs to insert its header in the packet buffer, in front of the application payload. Towards this end we need to:


1. **Allocate** space for the header → call `packetbuf_hdralloc(size)`
2. **Insert** the header in the packet buffer. **How?**



# Forwarding at a relay node – TODO 6

When a **non-sink node** receives a data packet, it should **forward** it to its **parent in unicast**:

1. Extract the collection header from the received payload
2. Increment the hop count in the header
3. Put the *updated header back* in the packet buffer
4. Send the packet to its current parent using unicast





```
struct collect_header {  
    linkaddr_t source;  
    uint8_t hops;  
} __attribute__((packed));
```

Set to 0 at the source. It must be incremented by 1 every time the packet is relayed

# Forwarding at the sink – TODO 6

When the **sink node** receives a data packet, it should (i) communicate the packet's source and hop count, and (ii) **deliver the app payload to the application**:


1. Extract the collection header from the received payload
2. Keep trace of the header's fields (source, hops) and update the hop count (hops + 1)
3. Remove the header from the payload
4. Call the *application* recv callback (recv\_cb) to inform the application about the received data (app payload, source, hops)




What is the “received payload”? How can we “extract” and/or “remove” our collection header from it?

## Extracting a header – TODO 6 (nodes & sink)

When we receive a packet from the underlying layer (`uc_recv()`), its payload starts at `packetbuf_dataptr()`



*Q<sub>1</sub>* What is such "**payload**"? It consists of both our **collection header** and the **application payload**.



*Q<sub>2</sub>* How can we **extract** the **collection header** from the received payload? Let's just use the `packetbuf_dataptr()` pointer to `memcpy` the header!

## Removing a header – TODO 6 (sink only)



**Q<sub>3</sub>** How can we **remove** the **collection header** from the received payload? Let's use **packetbuf\_hdrreduce(size)**


→ Shift the `packetbuf_dataptr()` pointer to the right of “size” bytes, thus making it point to the beginning of the app payload



**N.B.:** If the *application recv* callback is called after having removed the collection header, by looking at `packetbuf_dataptr()` the application can directly get the app payload!

# Packetbuf Management (SENDING)


**Application:** (i) clears the packetbuf, (ii) copies the data to the packetbuf, and (iii) sets the data length



**NB:** This is the packetbuf status when `my_collect_send(..)` is called!

# Packetbuf Management (SENDING)


**My Collect:** (i) *allocates space for the header (shifting the app data to the right → packetbuf\_hdralloc)* and  
(ii) *copies its header to the packet buffer*



**NB:** This is the packetbuf status when `unicast_send()` is called!


# Packetbuf Management (SENDING)

**Rime:** (i) *allocates* space for its own header (shifting the data to the right) and (ii) *copies* the header to packetbuf




# Packetbuf Management (SENDING)

**MAC:** (i) *allocates* space for its own header (shifting the data to the right) and (ii) *copies* the header to packetbuf




# Packetbuf Management (RECEPTION)

**RADIO:** (i) clears packetbuf, (ii) copies received packet to packetbuf, (iii) sets the data length (MAC + RIME + MC HDR + PLD), and (iv) calls the upper layer (MAC)




# Packetbuf Management (RECEPTION)

**MAC:** (i) *parses* MAC header, (ii) *reduces the header* (shifting the data pointer to the right of MAC HDR bytes), and (iii) *calls* the upper layer (RIME)



# Packetbuf Management (RECEPTION)


**Rime:** (i) *parses Rime header attributes*, (ii) *reduces header* (shifting the data pointer to the right of RIME HDR bytes), and (iii) *calls the upper layer (MY\_COLLECT)*



**N.B.:** This is the packetbuf status when uc\_recv () is called!

# Packetbuf Management (RECEPTION)


**My Collect—Non-sink:** (i) *parses* and *updates* collect header,  
(ii) *copies* the new header to packetbuf,  
(iii) *calls* the lower layer (RIME) via  
unicast\_send()



**N.B.: no need to shift it to the right here,  
we are not going to call the application!**

# Packetbuf Management (RECEPTION)


**My Collect—Sink:** (i) *parses* and updates collect header,  
(ii) *reduces* header (shifting the data pointer to  
the right of MC HDR bytes), and (iii) *calls* the  
upper layer (application, via `recv_cb`) .



**N.B.:** When the application `recv` callback is called `packetbuf_dataptr()` points to the *beginning* of the app data. The application can thus directly use this pointer to access the app payload!

# Packetbuf Management (RECEPTION)

**Application:** copies payload to a local buffer using  
`packetbuf_dataptr()` and processes the data



# Code template

## Download and unzip the provided code

- \$ unzip Lab7-exercise.zip

## Go to the code directory

- \$ cd Lab7-exercise/data-collection-template-Lab7



## If you have **not yet started** TODO 1 [Lab 6]

- Directly rely on this new code template!


## Otherwise:

- EITHER replace the `my_collect.c` file in this template with the one you developed in Lab 6, relying on the Lab 7 code template for the rest.
- OR keep working on the Lab 6 template, but (i) set `DATA_FORWARDING = 1` in `app.c`, and (ii) add the `cooja_no_gui` and `cooja_with_gui` folders and `parse-stats.py` to that template.

# Test your solution and ...

## Step 1 — Check that things work as expected

- Try the different Cooja simulation files in the `cooja_with_gui` folder
  - **Tips:** leverage (i) the Cooja timeline, and (ii) the filter functionality of the mote output tool



Check who TX and RX  
a given packet

The screenshot shows the "Mote output" window with a table of log entries. The columns are Time, Mote, and Message. The entries show a sequence of messages between mote ID 5 and ID 1. A blue highlight covers the row where mote 1 receives a message from mote 5. A yellow box at the bottom contains a filter command: "Filter: ID:5\\$+App:|ID:1\\$+App: Recv from 05".

Time	Mote	Message
00:00.996	ID:5	App: I am a normal node 05:00
00:34.000	ID:5	App: Send seqn 0
00:34.003	ID:5	App: Message seqn 0 sent correctly!
00:34.017	ID:1	App: Recv from 05:00 seqn 0 hops 2
01:01.289	ID:5	App: Send seqn 1
01:01.293	ID:5	App: Message seqn 1 sent correctly!
01:01.306	ID:1	App: Recv from 05:00 seqn 1 hops 2
01:35.149	ID:5	App: Send seqn 2
01:35.152	ID:5	App: Message seqn 2 sent correctly!

Visualise only the  
desired output.  
**NB: You can use regular  
expressions here!**

# Test your solution and ... Compete!

## Step 1 — Check that things work as expected


- Try the different Cooja simulation files in the `cooja_with_gui` folder
  - **Tips:** leverage (i) the Cooja timeline, and (ii) the filter functionality of the mote output tool

## Step 2 — Evaluate the protocol performance and compete with your classmates!

- Run a reasonably long simulation experiment with Cooja:  
`$ cooja testname.csc`
- After 15 simulated minutes, **save a log file** with your results



# How to save a Cooja log file



# Test your solution and ... Compete!

## Step 1 — Check that things work as expected

- Try the different Cooja simulation files in the `cooja_with_gui` folder
  - **Tips:** leverage (i) the Cooja timeline, and (ii) the filter functionality of the mote output tool

## Step 2 — Compare your protocol performance with your classmates!

- Run a longer simulation experiment with Cooja:  
`$ cooja testname.csc`
- After 15 simulated minutes, **save a log file** with your results
- Analyze the reliability of your protocol with:  
`$ python3 parse-stats.py my_log_file.log`
- Try to improve your solution!



# Let's automate: Cooja simulations without GUI!

Once your protocol is working reliably, you can run Cooja **without GUI**, increasing the simulation speed and making it easier to automate experiments.

## Steps:

- Emulate the Cooja simulation files provided in `cooja_no_gui`:  
`$ cooja_nogui testname_nogui.csc`
- By default:
  - 30 (simulated) minutes experiment
  - Motes output saved in `testname_nogui.log` file
  - Radio-on time saved in `testname_nogui_dc.log` file
- Analyze the reliability of your protocol with:  
`$ python3 parse-stats.py testname_nogui.log`

# Preparing Cooja simulations without GUI

To prepare other Cooja simulations without GUI, you can just copy and paste in your \*.csc files the following code snippet (from every `testname_nogui.csc` files we provide):

```
<plugin>
    org.contikios.coja.plugins.ScriptRunner
    <plugin_config>
        <script>SIM_SETTLING_TIME = 1000
                TIMEOUT(1800000);
                .
                .
                .
        </script>
        <active>true</active>
    </plugin_config>
</plugin>
```

You can change the experiment duration by modifying this value!

**NOTE: the timeout is in ms**

# Expected performance of your protocol

Once you have implemented all the basic functionalities we have discussed, your system should achieve *similar performance*

```
MacBook-Air-di-Matteo:exp matteo$ python ./parse-stats.py test_nogui.log
```

```
Namespace(logfile='test_nogui.log', testbed=False)
```

```
LogFile: test_nogui.log
```

```
Cooja simulation
```

```
##### Node Statistics #####
```

```
Node 2: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 3: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 4: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 5: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 6: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 7: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 8: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 9: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 10: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
##### Overall Statistics #####
```

```
Total Number of Packets Sent: 531
```

```
Total Number of Packets Received: 531
```

```
Overall PDR = 100.00%
```

```
Overall PLR = 0.00%
```

```
MacBook-Air-di-Matteo:exp matteo$ python ./parse-stats.py test_more_random_nogui.log
```

```
Namespace(logfile='test_more_random_nogui.log', testbed=False)
```

```
LogFile: test_more_random_nogui.log
```

```
Cooja simulation
```

```
##### Node Statistics #####
```

```
Node 2: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 3: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 4: TX Packets = 59, RX Packets = 58, PDR = 98.31%, PLR = 1.69%
```

```
Node 5: TX Packets = 59, RX Packets = 58, PDR = 98.31%, PLR = 1.69%
```

```
Node 6: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 7: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 8: TX Packets = 59, RX Packets = 58, PDR = 98.31%, PLR = 1.69%
```

```
Node 9: TX Packets = 59, RX Packets = 59, PDR = 100.00%, PLR = 0.00%
```

```
Node 10: TX Packets = 59, RX Packets = 58, PDR = 98.31%, PLR = 1.69%
```

```
##### Overall Statistics #####
```

```
Total Number of Packets Sent: 531
```

```
Total Number of Packets Received: 527
```

```
Overall PDR = 99.25%
```


```
Overall PLR = 0.75%
```

There is still room for improvements... Can you push it further?

Does it work?  
Try it in CLOVES  
with Firefly nodes!

# Suggested topology to start with

**DISI POVOII floor (part of DEPT) → 36 nodes, use node 1 as sink**



Other topologies are available! Check them at

<https://research.iottestbed.disi.unitn.it/maps/>

**N.B.: Before to use different topologies, remember to update:**

1. `linkaddr_t sink` in `app.c`
2. `sink_id, addr_id_map` in `parse-stats.py`

# Running a test in CLOVES

## Steps:

1. **Compile your code for the Zolertia Firefly platform:**

```
$ make TARGET=zoul
```

2. **Log into the CLOVES interface:**

Start at `research.iottestbed.disi.unitn.it`, and log in with your credentials. Then navigate to “Create job”.

3. **Prepare a job:**

The screenshot shows the 'Create job' interface in the CLOVES web application. It consists of two main sections: 'Timeslot info' on the left and 'Binary file 1' on the right.

**Timeslot info:**

- Island: DEPT
- Start time: ASAP
- Start time date: 13/11/2025 15:40
- Duration: 320

**Binary file 1:**

- Upload file: Enabled (blue toggle switch)
- Hardware: firefly
- Bin file: app.bin
- Sfoglia... (Browse button)
- Targets: disi\_povo2
- Programaddress: 0x00200000

4. **Wait for the job to complete & download the logs:**

Check (and download) completed jobs at:

[https://research.iottestbed.disi.unitn.it/jobs\\_completed/](https://research.iottestbed.disi.unitn.it/jobs_completed/)

5. **Parse and analyze the logs:**

```
$ python3 parse-stats.py your_path/job.log --testbed
```

# Performance at night

```
Lab6_7-solution/data-collection-template/cloves/job_44552 » python ../../parse-stats.py job.log --testbed  
Namespace(logfile='job.log', testbed=True)  
LogFile: job.log  
Testbed experiment
```

```
##### Node Statistics #####
```

```
Node 2: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 3: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 4: TX Packets = 16, RX Packets = 16, PDR = 100.00%, PLR = 0.00%  
Node 5: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 6: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 7: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 8: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 9: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 10: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 11: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 12: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 13: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 14: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 15: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 16: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 17: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 18: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 19: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 20: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 21: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 22: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 23: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 24: TX Packets = 16, RX Packets = 16, PDR = 100.00%, PLR = 0.00%  
Node 25: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 26: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 27: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 28: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 29: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 30: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 31: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 32: TX Packets = 16, RX Packets = 16, PDR = 100.00%, PLR = 0.00%  
Node 33: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 34: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 35: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%  
Node 36: TX Packets = 15, RX Packets = 15, PDR = 100.00%, PLR = 0.00%
```

```
##### Overall Statistics #####
```

```
Total Number of Packets Sent: 528  
Total Number of Packets Received: 528  
Overall PDR = 100.00%  
Overall PLR = 0.00%
```

## N.B.:

- Night-time experiment on *channel 25!* Results might **highly vary** in function of the amount of interference present in the channel!
- Be careful, during the day the amount of interference is typically higher!
- Try to run an experiment during the day on channel 11 (typically highly exposed to Wi-Fi traffic); do you notice any difference?

What about  
energy consumption?


# NullRDC Vs. ContikiMAC

So far, we used **NullRDC** → Radio is *always ON, high power consumption!*



More energy efficient radio duty-cycle (RDC) layers are available!  
Check them at [contiki-uwb/contiki/core/net/mac](https://contiki-uwb.github.io/contiki/core/net/mac)

- Try **ContikiMAC\*** and check how it affects performance



\* A. Dunkels, "The ContikiMAC Radio Duty Cycling Protocol", *SICS Technical Report*, 2011

# NullRDC Vs. ContikiMAC

So far, we used **NullRDC** → Radio is ***always ON, high power consumption!***



More energy efficient radio duty-cycle (RDC) layers are available!  
Check them in core/net/mac

- Try **ContikiMAC\*** and check how it affects performance

## Steps:

1. Change the RDC protocol in your `project-conf.h` file:

```
/* #define NETSTACK_RDC nullrdc_driver */  
#define NETSTACK_RDC contikimac_driver
```

2. Check the PDR of your system, do you notice any difference?
3. Compare ContikiMAC and NullRDC radio-on time:
  - Cooja Timeline: Can you recognize the peculiar ContikiMAC radio on-off pattern?
  - Radio duty-cycle: Cooja -> Tools -> Mote radio duty cycle